



Open-Source Hardware for Heterogeneous Computing with ESP and RISC-V

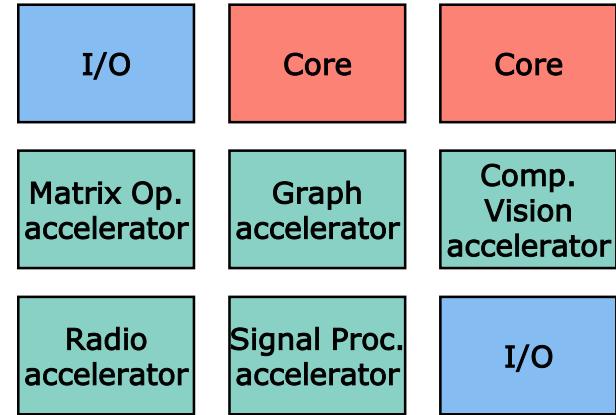
Luca P. Carloni

Spring 2022 RISC-V Week
Paris - May 3, 2022



The Age of Heterogeneous Computing

- State-of-the-art SoC architectures integrate increasingly diverse sets of components
 - different CPUs, GPUs, hardware accelerators, memory hierarchies, I/O peripherals, sensors, reconfigurable engines, analog blocks...
- The migration towards heterogeneous SoC architectures will accelerate, across almost all computing domains
 - IoT devices, mobile devices, embedded systems, automotive electronics, avionics, data centers and even supercomputers
- The set of heterogeneous SoCs in production in any given year will be itself heterogeneous!
 - no single SoC architecture will dominate all the markets!



Heterogeneity Increases Design Complexity

- Heterogeneous architectures produce higher energy-efficient performance, but make more difficult the tasks of design, verification and programming
 - at design time, diminished regularity in the system structure, chip layout
 - at runtime, more complex hardware/software and management of shared resources
- With each SoC generation, the addition of new capabilities is increasingly limited by engineering effort and team sizes
- The biggest challenges are (and will increasingly be) found in the **complexity of system integration**



Open-Source Hardware (OSH)

- An opportunity to reenergize the innovation in the semiconductor and electronic design automation industries
- The OSH community is gaining momentum
 - many diverse contributions from both academia and industry
 - multi-institution organizations
 - government programs



RISC-V: The Free and Open RISC
Instruction Set Architecture

NVDLA.org



OpenPiton

Image Sources:

<https://chipsalliance.org/>

<https://github.com/nvdla>

<https://www.openhwgroup.org/>

<https://parallel.princeton.edu/openpiton/>

<https://pulp-platform.org/>

<https://riscv.org/>

The Open Challenge of Open-Source Hardware

- To date, however, most OSH projects are focused on the development of individual SoC components, such as a processor core or an accelerator
- This leaves open a critical challenge:

How can we realize a complete SoC for a given target application domain by efficiently reusing and combining a variety of independently developed, heterogeneous, OSH components, especially if these components are designed by separate organizations for separate purposes?



The Concept of Platform

- Innovation in SoC architectures and their design methodologies is needed to promote design reuse and collaboration
 - Architectures and methodologies must be developed together
- *Platform = architecture + methodology*
 - An SoC architecture enables design reuse when it simplifies the integration of many components that are independently developed
 - An SoC methodology enables design collaboration when it allows designers to choose the preferred specification languages and design flows for the various components
- An effective combination of architecture and methodology is a platform that maximizes the potential of open-source hardware
 - by scaling-up the number of components that can be integrated in an SoC and by enhancing the productivity of the designers who develop and use them



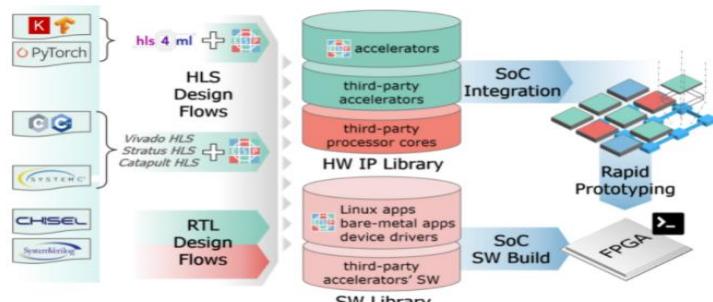
ESP : An Open-Source Platform for SoC Design

The screenshot shows the top navigation bar with links: Home, Release, Resources (with a dropdown), News, Press, Team, Contact. Below the navigation is the ESP logo: "ESP" in blue and orange, followed by "the open-source SoC platform". On the right side of the header is the URL "www.esp.cs.columbia.edu".



The ESP Vision

ESP is an open-source research platform for heterogeneous system-on-chip design that combines a scalable tile-based architecture and a flexible system-level design methodology.

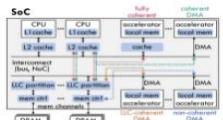


ESP provides three accelerator flows: RTL, high-level synthesis (HLS), machine learning frameworks. All three design flows converge to the ESP automated SoC integration flow that generates the necessary hardware and software interfaces to rapidly enable full-system prototyping on FPGA.

Overview

A video player interface with the title "Quick Intro to ESP - The Open So...". Below the video player is the text "The Open Source SoC Platform".

Latest Posts



Paper accepted at MICRO 2021

Our paper "Cohmeleon: Learning-Based Orchestration of Accelerator Coherence in Heterogeneous SoCs" will be presented in October at MICRO 2021.

[Read more](#)

Published: Sep 14, 2021

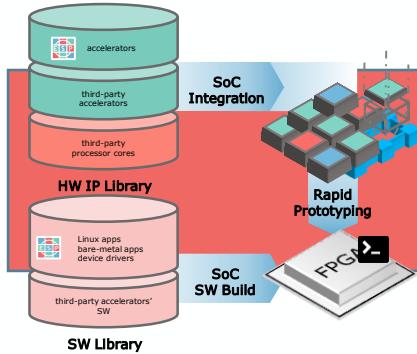
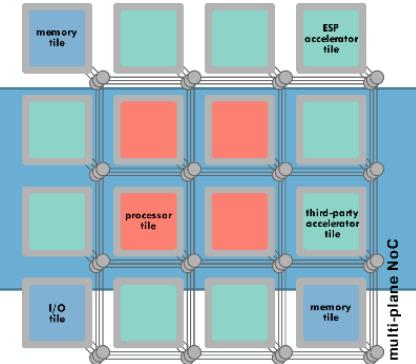


Paper published in



Outline

The ESP Architecture



The ESP Methodology

Research & Teaching with ESP

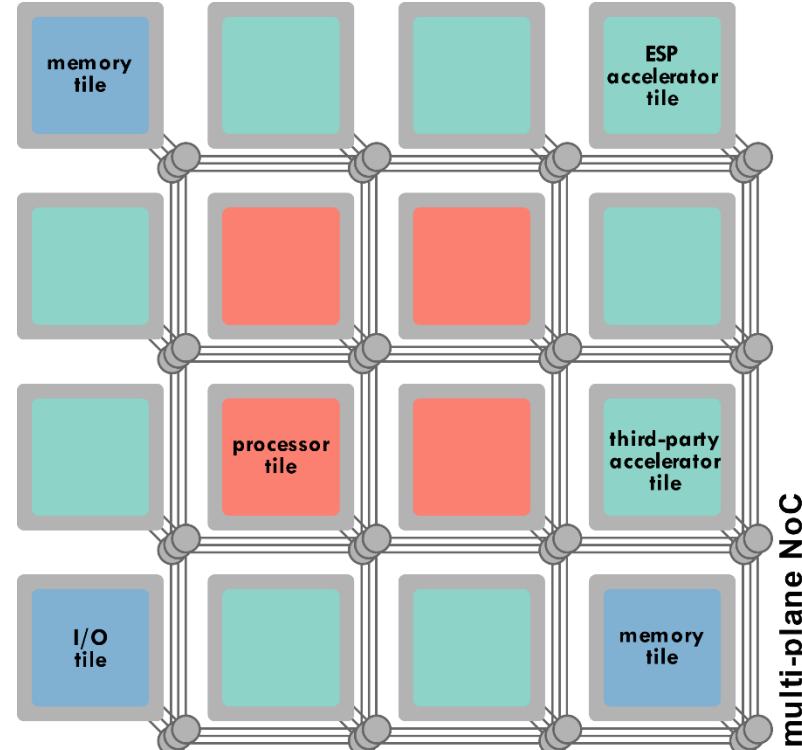
The screenshot shows the ESP website homepage. At the top, it says "ESP the open-source SoC platform". Below that is a "The ESP Vision" section with a brief description and a diagram of the design flow. To the right is a "Latest Posts" sidebar featuring a thumbnail of a desert landscape and a link to an upcoming talk by Professor Carloni at VLSSoC 2020.



ESP Architecture

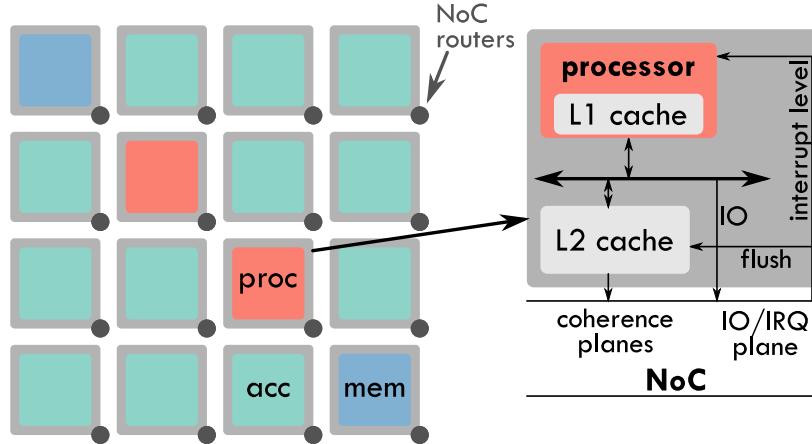
- RISC-V Processors
- Many-Accelerator
- Distributed Memory
- Multi-Plane NoC

The ESP architecture implements a **distributed** system, which is **scalable**, **modular** and **heterogeneous**, giving processors and accelerators similar weight in the SoC



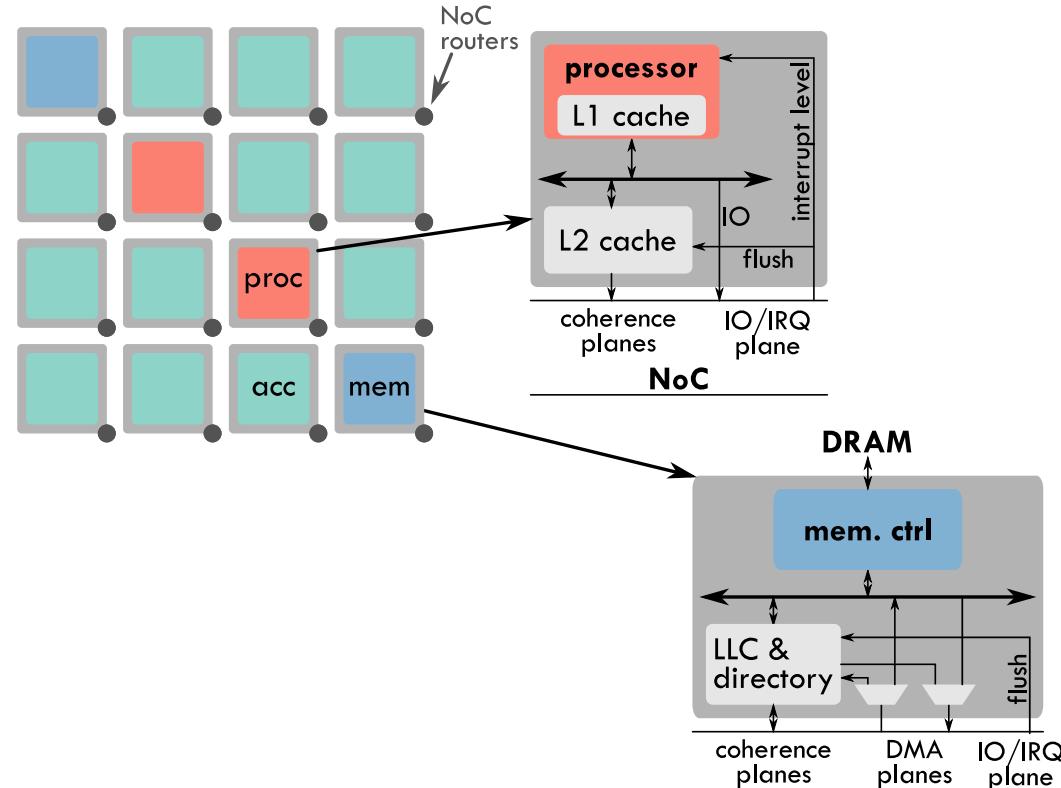
ESP Architecture: Processor Tile

- Processor off-the-shelf
 - RISC-V CVA6-Ariane (64 bit)
 - SPARC V8 Leon3 (32 bit)
 - RISC-V IBEX (32 bit)
 - L1 private cache
- L2 private cache
 - Configurable size
 - MESI protocol
- IO/IRQ channel
 - Un-cached
 - Accelerator config. registers, interrupts, flush, UART, ...



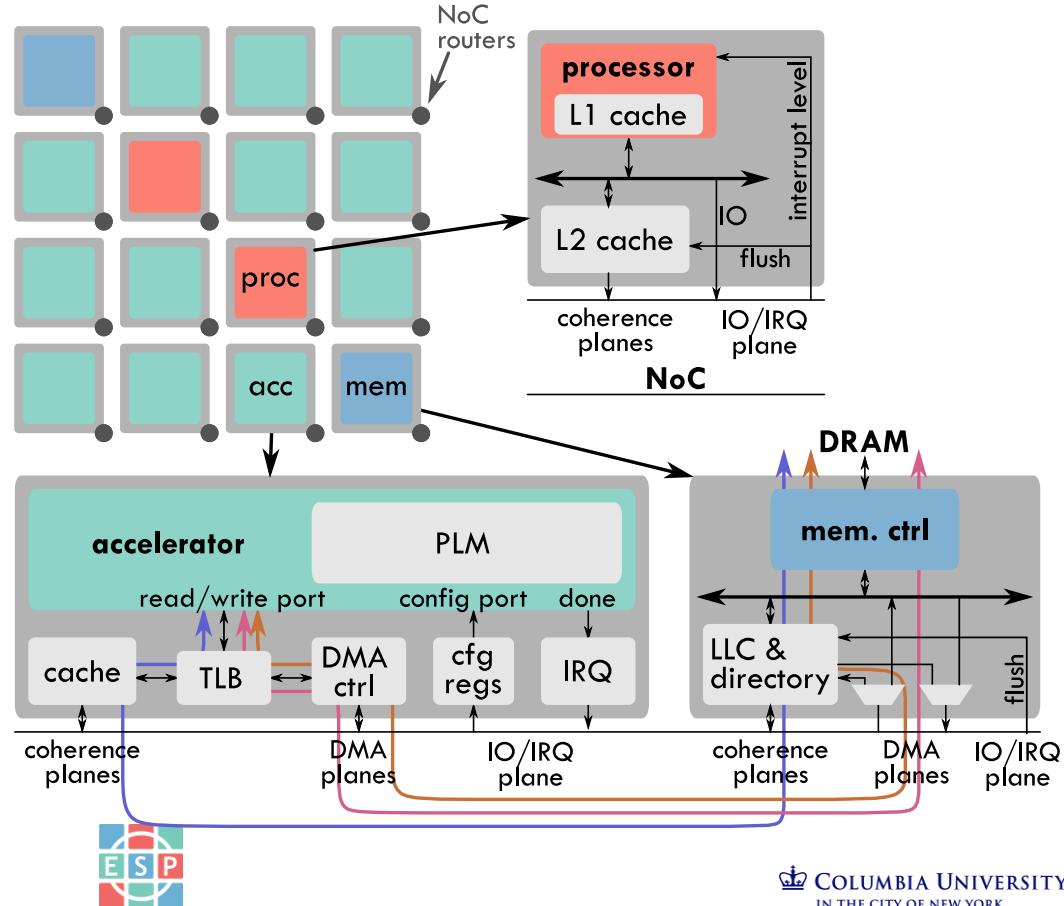
ESP Architecture: Memory Tile

- External Memory Channel
- LLC and directory partition
 - Configurable size
 - Extended MESI protocol
 - Supports coherent-DMA for accelerators
- DMA channels
- IO/IRQ channel

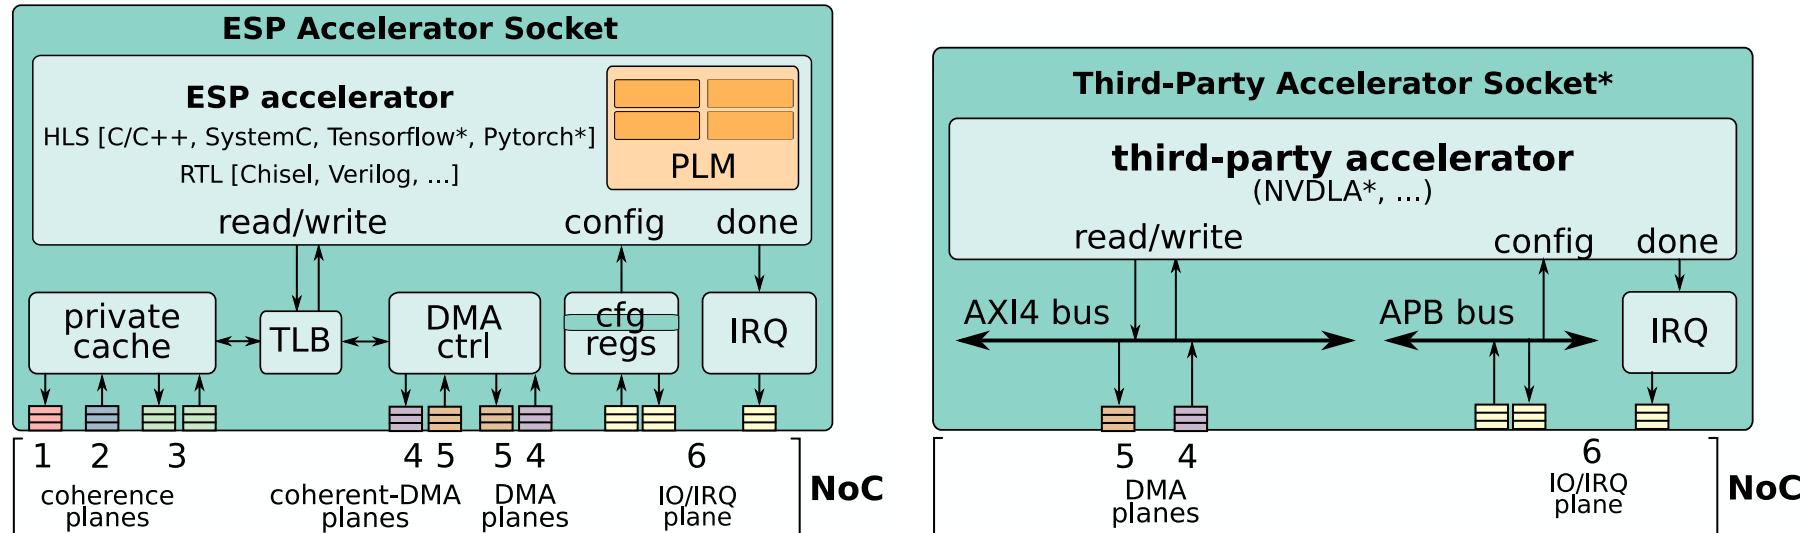


ESP Architecture: Accelerator Tile

- Accelerator Socket w/ Platform Services
 - Direct-memory-access
 - Run-time selection of coherence model:
 - Fully coherent
 - LLC coherent
 - Non coherent
 - User-defined registers
 - Distributed interrupt



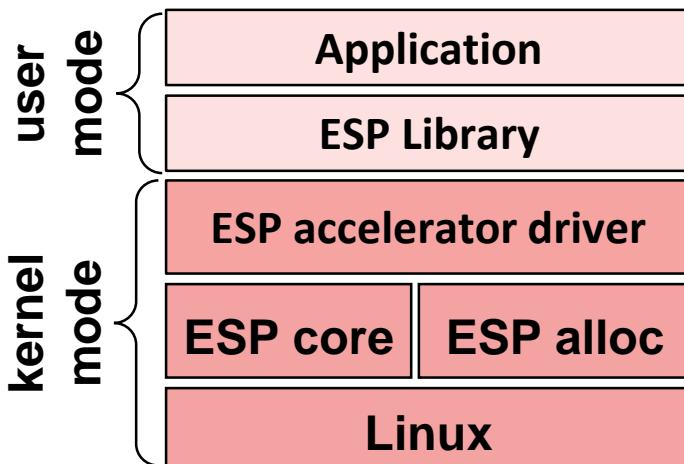
ESP Accelerator Socket



ESP Software Socket

- **ESP accelerator API**

- Generation of device driver and unit-test application
- Seamless shared memory



```
/*
 * Example of existing C application with ESP
 * accelerators that replace software kernels 2, 3,
 * and 5. The cfg_k# contains buffer and the
 * accelerator configuration.
 */
{
    int *buffer = esp_alloc(size);

    for (...) {
        kernel_1(buffer,...); /* existing software */

        esp_run(cfg_k2);      /* run accelerator(s) */
        esp_run(cfg_k3);

        kernel_4(buffer,...); /* existing software */
        esp_run(cfg_k5);
    }

    validate(buffer);          /* existing checks */

    esp_free();                /* memory free */
}
```



ESP Platform Services

Accelerator tile

DMA

Reconfigurable coherence

Point-to-point

ESP or AXI interface

DVFS controller

Processor Tile

Coherence

I/O and un-cached memory

Distributed interrupts

DVFS controller

Miscellaneous Tile

Debug interface

Performance counters access

Coherent DMA

Shared peripherals (UART, ETH, ...)

Memory Tile

Independent DDR Channel

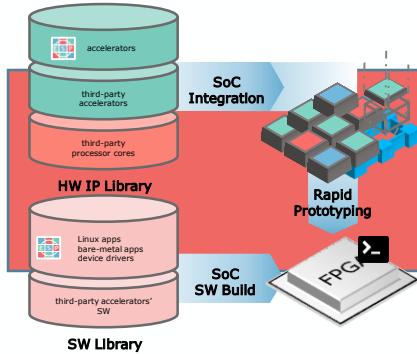
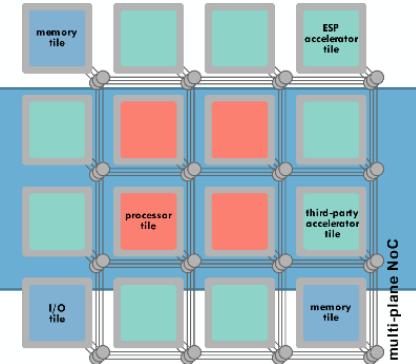
LLC Slice

DMA Handler



Outline

The ESP Architecture



The ESP Methodology

Research & Teaching with ESP

The ESP website homepage features a navigation bar with links to Home, Resources, News, Press, Team, and Contact. The main content area is titled "The ESP Vision" and describes ESP as an open-source research platform for heterogeneous system-on-chip design. It highlights the "HLS Design Flows" (using PyTorch, HLS 4, and Vivado HLS), "RTL Design Flows" (using SystemC, C++, and Chisel), and "SoC Integration" tools. A sidebar on the right shows "Latest Posts" about an upcoming talk at VLSSoC 2020 and includes a "Read more" button and a "Published Sep 11, 2020" timestamp.

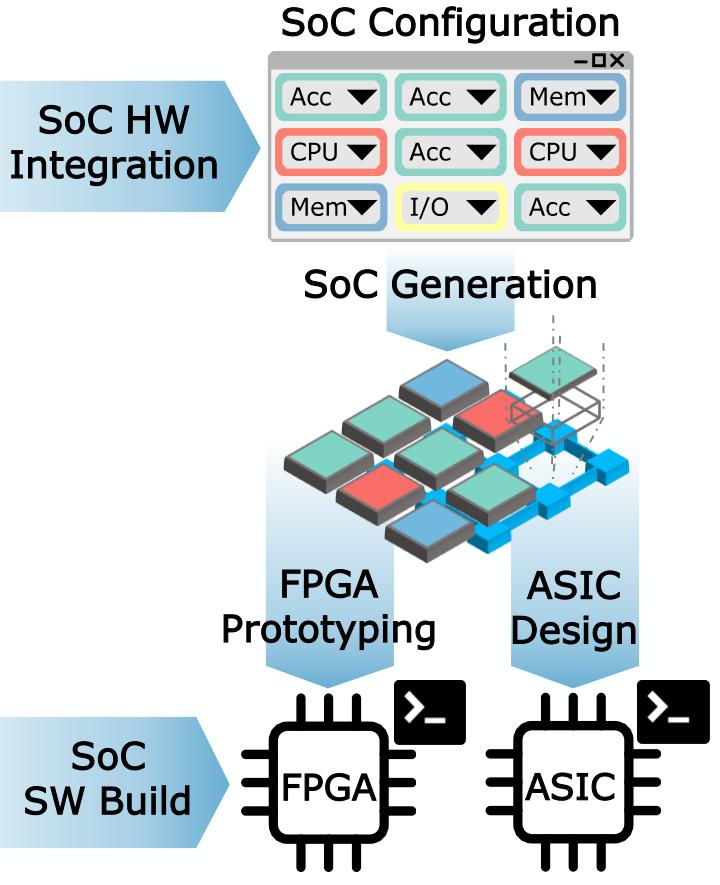
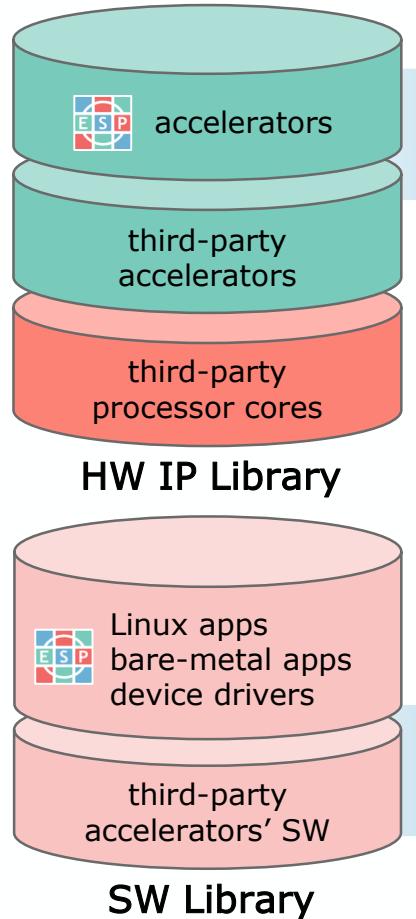
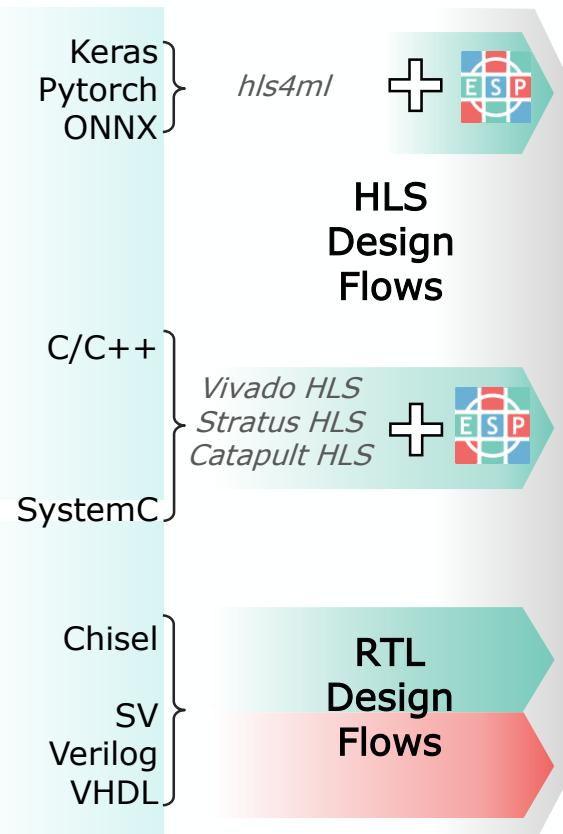


The Pillars of the ESP Approach

- Develop platforms, not just architectures
 - A platform combines an architecture and a companion design methodology
- Move from a processor-centric to an SoC-centric perspective
 - The processor core is just one component among many others
- Raise the level of abstraction
 - Move from RTL design to domain-specific *system-level design* with high-level synthesis...
 - ...but keep supporting different abstraction levels and design flows
- Promote Open-Source Hardware
 - Build libraries of reusable components
 - Support the integration of third-party IP components



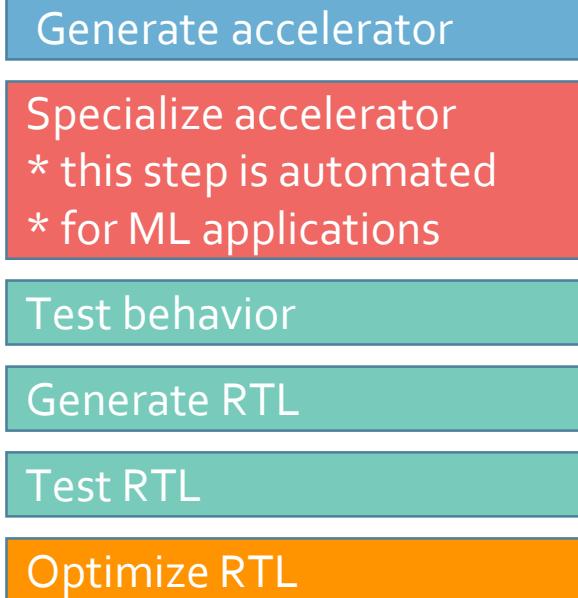
The ESP Vision: Domain Experts Can Design SoCs



ESP Methodology In Practice



Accelerator Flow

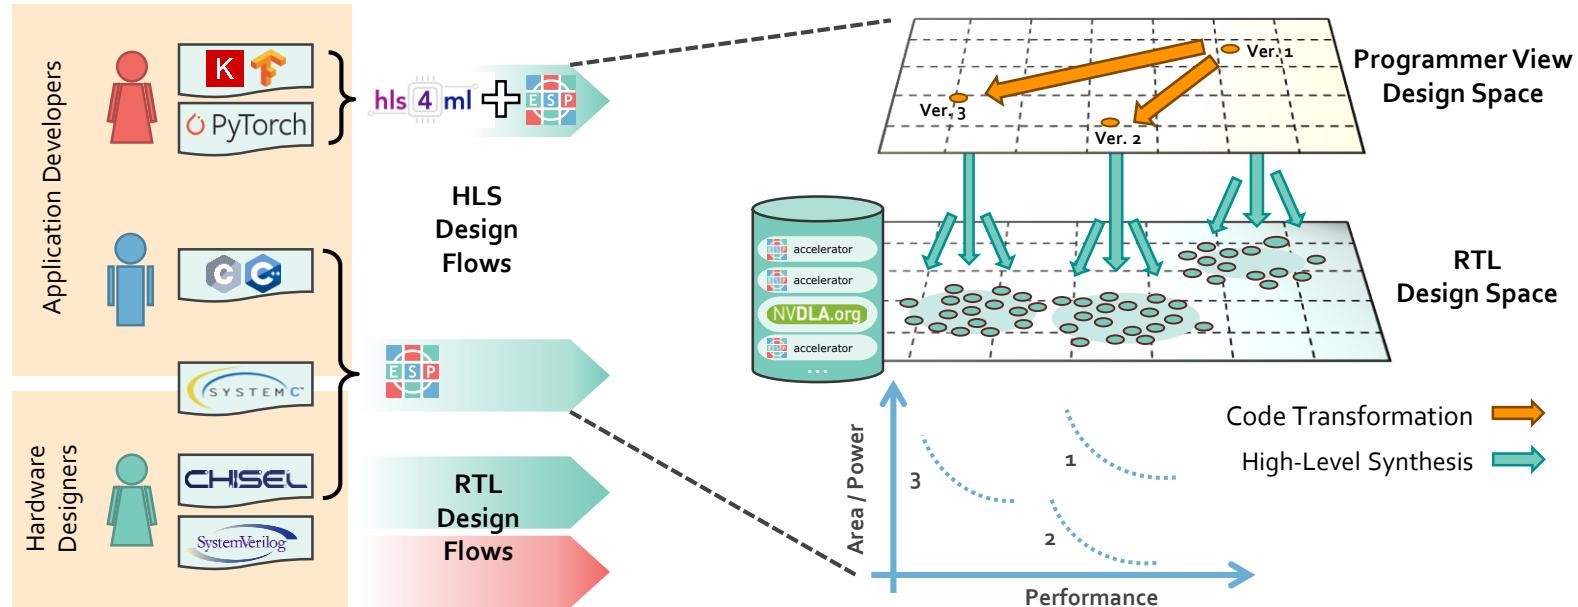


SoC Flow



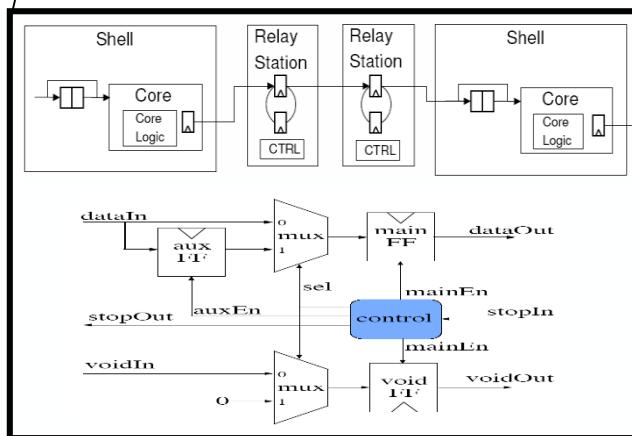
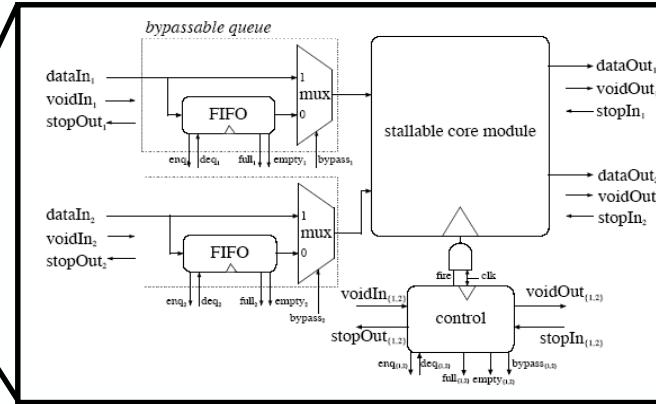
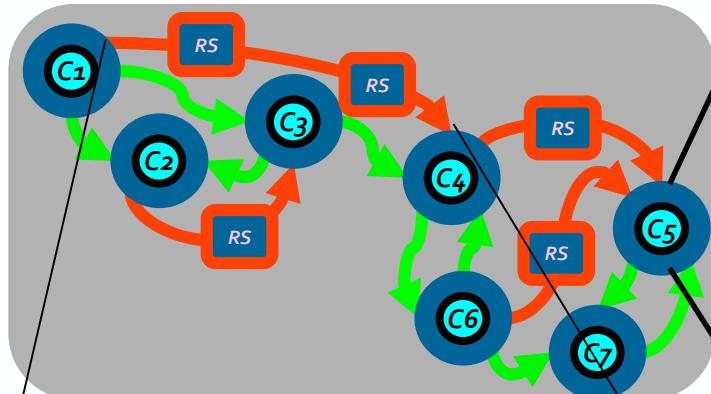
ESP Accelerator Flow

Developers focus on the **high-level specification**, decoupled from memory access, system communication, hardware/software interface



Retrospective: Latency-Insensitive Design

[Carloni et al. '99]



Latency-Insensitive Design

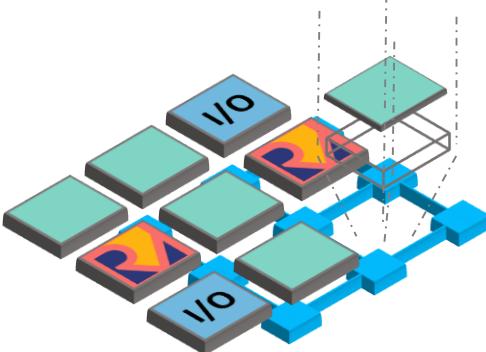
- is the foundation for the *flexible synthesizable RTL representation*
- anticipates the separation of computation from communication that is proper of TLM with SystemC
 - through the introduction of the Protocols & Shell paradigm



ESP Interactive Flow for SoC Integration



SoC Integration



The screenshot shows the "ESP SoC Generator" application window. The interface is divided into several sections:

- General SoC configuration:**
 - virtexup
 - ETH FPNw
 - No JTAG
 - Eth (192.168.1.2)
 - Use SGMII
 - No SVGA
 - With synchronizers
- Data transfers:**
 - Bigphysical area
 - Scatter/Gather
- Cache Configuration:**
 - Cache En.:
 - L2 SETS: 512
 - L2 WAYS: 4
 - LLC SETS: 1024
 - LLC WAYS: 16
 - ACC L2 SETS: 512
 - ACC L2 WAYS: 4
- CPU Architecture:**
 - Core: ariane
- NoC configuration:**
 - Rows: 2 Cols: 2
 - Config
 - Monitor DDR bandwidth
 - Monitor memory access
 - Monitor injection rate
 - Monitor router ports
 - Monitor accelerator status
 - Monitor L2 Hit/Miss
 - Monitor LLC Hit/Miss
 - Monitor DVFS
- NoC Tile Configuration:**

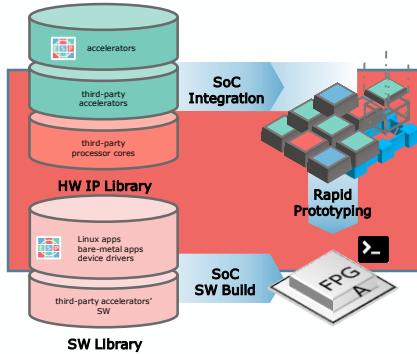
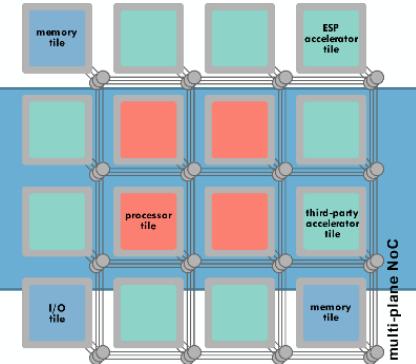
(0,0)	(0,1)
mem	cpu
mem	cpu
(1,0)	(1,1)
empty	IO
empty	IO

 - (0,0) mem
 - (0,1) cpu
 - (1,0) empty
 - (1,1) IO- Has L2 | Clk Reg: 0 | Has PLL | CLK BUF
- Has L2 | Clk Reg: 0 | Has PLL | CLK BUF
- Has L2 | Clk Reg: 0 | Has PLL | CLK BUF
- Has L2 | Clk Reg: 0 | Has PLL | CLK BUF

- Bottom Panel:**
- Generate SoC config


Outline

The ESP Architecture



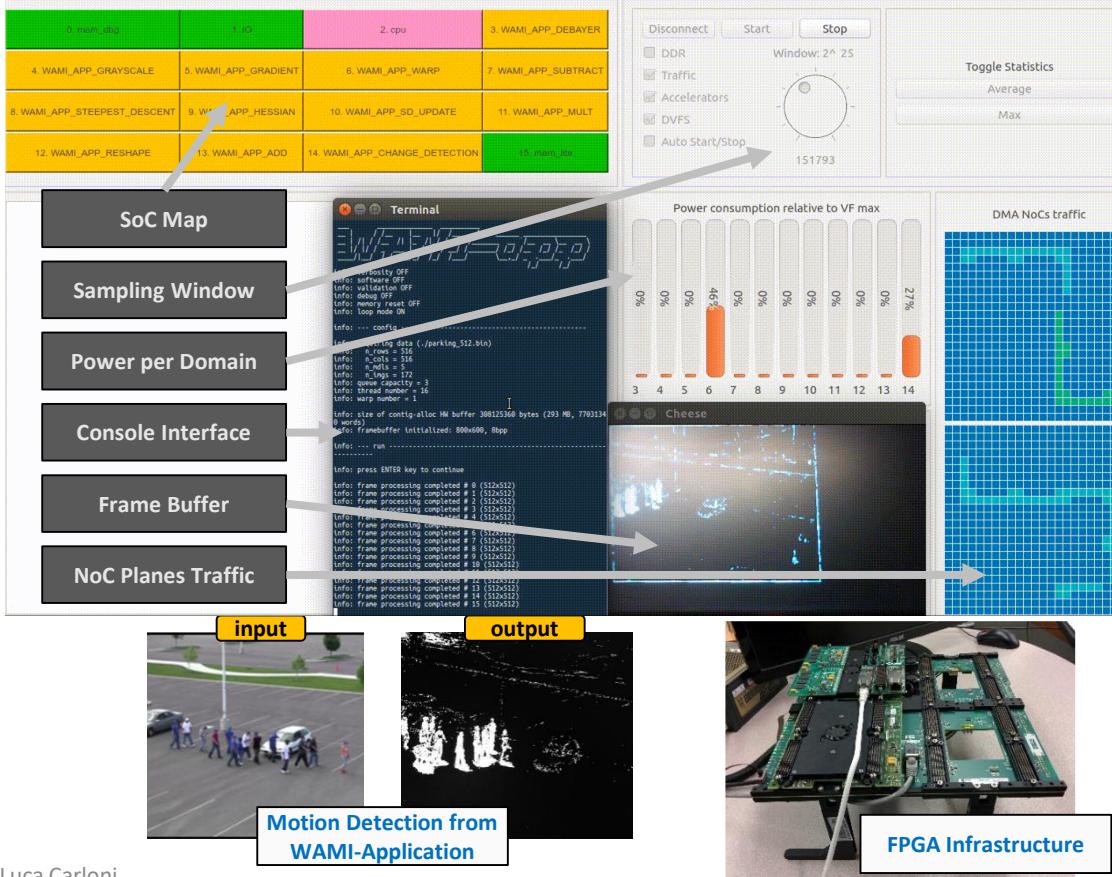
The ESP Methodology

Research & Teaching with ESP

The screenshot shows the ESP website homepage. At the top, it says "ESP the open-source SoC platform". Below that is a "The ESP Vision" section with a brief description and a diagram of the design flow. To the right is a "Latest Posts" sidebar featuring a thumbnail of a desert landscape and a link to an upcoming talk by Professor Carloni at VLSSoC 2020.



Example of a System We Built: FPGA Prototype to Accelerate Wide-Area Motion Imagery

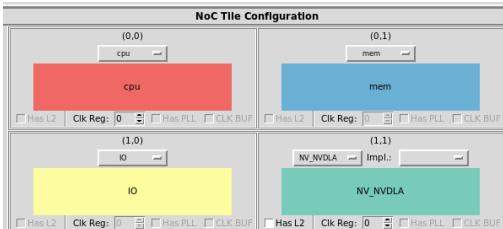


- **Design:** Complete design of WAMI-App running on an FPGA implementation of an ESP architecture
 - featuring 1 embedded processor, 12 accelerators, 1 five-plane NoC, and 2 DRAM controllers
 - SW application running on top of Linux while leveraging multi-threading library to program the accelerators and control their concurrent, pipelined execution
 - **Five-plane, 2D-mesh NoC** efficiently supports multiple independent frequency domains and a variety of platform services

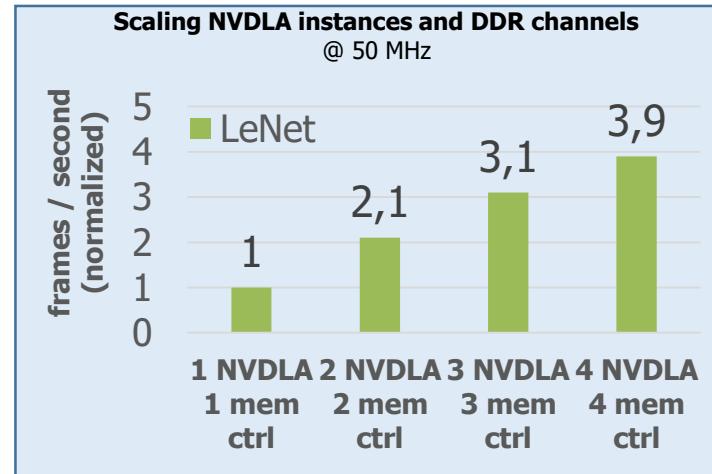
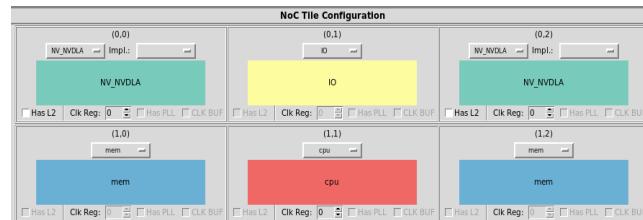
[P. Mantovani , L. P. Carloni et al., *An FPGA-Based Infrastructure for Fine-Grained DVFS Analysis in High-Performance Embedded Systems*, DAC 2016]

Seamless Integration of NVDLA Accelerators

- New design flow of general applicability to integrate third-party accelerators
 - demonstrated w/ NVIDIA Deep Learning Accelerator (NVDLA)
- Transparent accelerator integration
 - original software apps can run “as is”
- Linear performance scalability
 - when scaling up NVDLA instances with DDR channels



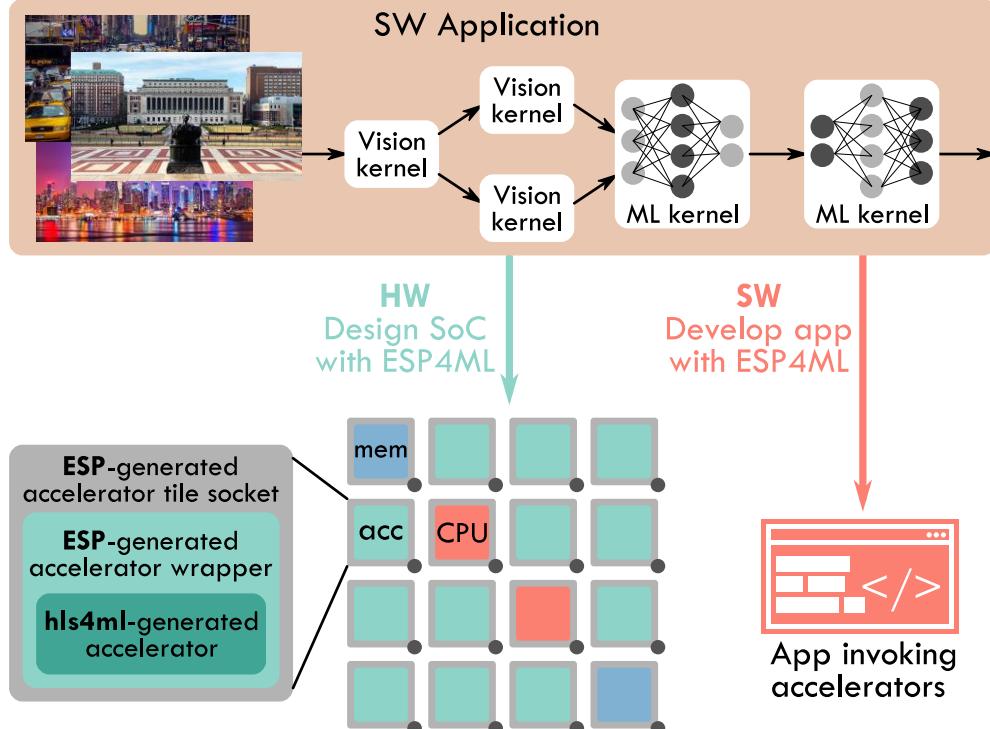
[D. Giri et al. “Ariane + NVDLA: Seamless Third-Party IP Integration with ESP”, CARRV’20]



ESP4ML

Open-source design flow to build and program SoCs for ML applications

- Combines  and 
 - **ESP** is a platform for heterogeneous SoC design
 - **hls4ml** automatically generates accelerators from ML models
- Main contributions to ESP:
 - Automated integration of hls4ml accelerators
 - Accelerator-accelerator communication
 - Accelerator invocation API



[D. Giri, K.-L. Chiu, G. Di Guglielmo, P. Mantovani, and L. P. Carloni. "ESP4ML: Platform-Based Design of Systems-on-Chip for Embedded Machine Learning", DATE '20]



hls4ml

- Open-source tool developed by **Fast ML Lab**
- Translates ML algorithms into accelerator specifications that are synthesizable with high-level synthesis tools for both FPGA and ASIC implementations
- Born for high-energy physics (small and ultra-low latency networks), it is gaining broad applicability and a growing community of contributors and users

arXiv:2103.05579v3 [cs.LG] 23 Mar 2021

hls4ml: An Open-Source Codesign Workflow to Empower Scientific Low-Power Machine Learning Devices

Farah Fahim*
Benjamin Hawks
Christian Herwig
James Hirschauer
Sergo Jindariani
Nhan Tran*
Fermilab
Batavia, IL, USA

Manuel Blanco Valentin
Josiah Hester
Yingyi Luo
John Mamish
Seda Orgencili-Memik
Northwestern University
Evanston, IL, USA

Scott Hauck
Shih-Chieh Hsu
University of Washington
Seattle, WA, USA

Duc Hoang
Rhodes College
Memphis, TN, USA

Luca P. Carloni
Giuseppe Di Guglielmo
Columbia University
New York, NY, USA

Thea Arrestad
Hamza Javed
Vladimir Loncar
Maurizio Pierini
Adrian Alan Pol
Sioni Summers
European Organization for Nuclear
Research (CERN)
Geneva, Switzerland

Jennifer Ngadiuba
Caltech
Pasadena, CA, USA

Edward Kreinar
HawkEye360
Herndon, VA, USA

Philip Harris
Jeffrey Krupa
Dylan Rankin
MIT
Cambridge, MA, USA

Javier Duarte
UC San Diego
La Jolla, CA, USA
jduarte@ucsd.edu

Mia Liu
Purdue University
West Lafayette, IN, USA

Zhenbin Wu
University of Illinois at Chicago
Chicago, IL, USA

ABSTRACT

Accessible machine learning algorithms, software, and diagnostic tools for energy-efficient devices and systems are extremely valuable across a broad range of application domains. In scientific domains, real-time near-sensor processing can drastically improve experimental design and accelerate scientific discoveries. To support domain scientists, we have developed hls4ml, an open-source software-hardware codesign workflow to interpret and translate machine learning algorithms for implementation with both FPGA and ASIC technologies. In this paper, we describe the essential features of the hls4ml workflow including network optimization

*Also affiliated with Northwestern University

Permission is granted by the copyright holders for personal use or classroom use, without prior permission or fee. For other uses, such as redistribution, prior permission or fee is required. This material is presented to ensure timely dissemination of scholarly and technical work. Copyright for third-party components of this work must be honored. All other rights reserved. Contact the owner/authors.

TMW Research Symposium 21, March 2021, San Jose, CA
© 2021 Copyright held by the owner/authors.

techniques—such as pruning and quantization-aware training—which can be incorporated naturally into the device implementations. We expand on previous hls4ml work by extending capabilities and techniques toward low-power implementations and increased usability: new Python APIs, quantization-aware pruning, end-to-end FPGA workflows, long pipeline kernels for low power, and new device backends include an ASIC workflow. Taken together, these and continued efforts in hls4ml will arm a new generation of domain scientists with accessible, efficient, and powerful tools for machine-learning-accelerated discovery.

KEYWORDS

hls4ml, machine learning, neural networks, tinyML, FPGA, ASIC, low-power, low-latency

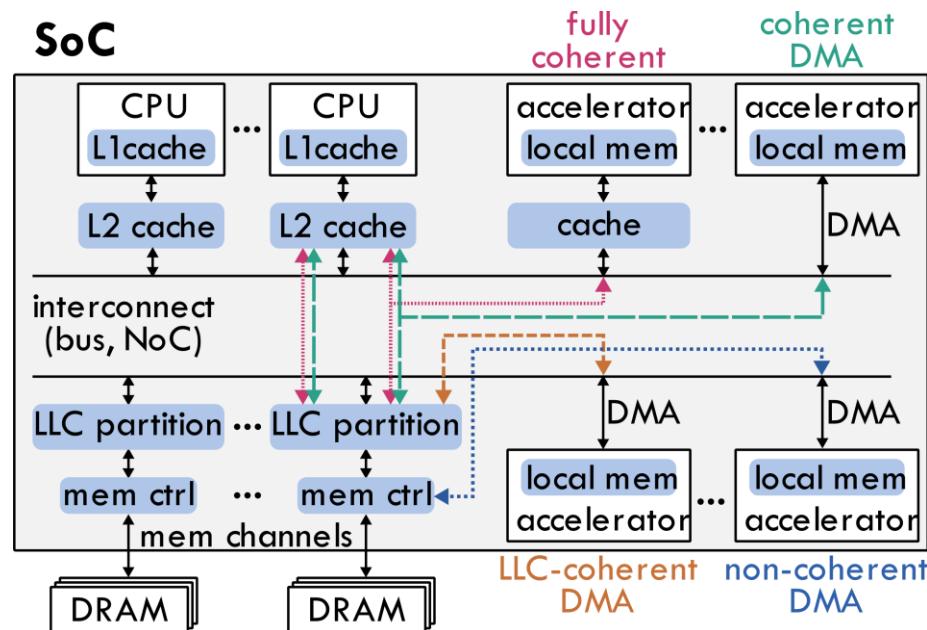
ACM Reference Format:

Farah Fahim, Benjamin Hawks, Christian Herwig, James Hirschauer, Sergo Jindariani, Nhan Tran, Luca P. Carloni, Giuseppe Di Guglielmo, Philip Harris, Jeffrey Krupa, Dylan Rankin, Manuel Blanco Valentin, Josiah Hester, Yingyi Luo, John Mamish, Seda Orgencili-Memik, Thea Arnestad, Hamza Javed, Vladimir Loncar, Maurizio Pierini, Adrian Alan Pol, Sioni Summers, Javier Duarte, Scott Hauck, Shih-Chieh Hsu, Jennifer Ngadiuba, Mia Liu, Duc



Cohmeleon: Learning-Based Orchestration of Accelerator Coherence in Heterogeneous SoCs

- Accelerator performance can vary greatly based on **coherence modes**
 - SoCs should support multiple coherence modes for optimal performance
- **Reinforcement learning** can be used to automatically manage coherence mode decisions
- With little overhead, Cohmeleon provides significant performance benefits for multiple objectives



[J. Zuckerman, D. Giri, J. Kwon, P. Mantovani, and L. P. Carloni **Cohmeleon: Learning-Based Orchestration of Accelerator Coherence in Heterogeneous SoCs**

IEEE/ACM International Symposium on Microarchitecture (MICRO-54), 2021.]

©Luca Carloni



Teaching with ESP

<https://www.esp.cs.columbia.edu>



Home Release Resources ▾ News Press Team Contact

Teaching

Class projects

1) Design and integration of an accelerator with ESP

For this project each student will use [ESP](#) to design one or more accelerators and to integrate them in a system-on-chip (SoC), capable of booting Linux. Then the student will evaluate the SoC both with RTL simulation and on FPGA.

To get a more practical sense of the project, you should familiarize yourself with ESP by using the resources on this website. Specifically:

- Check out the ESP website [Homepage](#) including the short introductory video.
- Watch the 16 minutes overview video in the [Documentation](#) section.
- Watch the videos and read the guides of the relevant hands-on tutorials available in the [Documentation](#) section. Especially relevant are the “*How to: setup*”, “*How to: design a single-core SoC*” guides and the “*How to: design an accelerator in ...*” guide that applies to your specific project.
- Explore the rest of the website to get the full picture of the ESP project.

Accelerator flows

For your project proposal, you are asked to choose which design flow you want to use to build your accelerator.

ESP offers multiple accelerator design flows: Stratus HLS flow (accelerator designed in SystemC), the Vivado HLS flow (accelerator designed in C/C++), the Catapult HLS flow (accelerator designed in C/C++) and the hls4ml flow (accelerators designed in Keras/Pytorch/ONNX).

Other options include designing the accelerator in RTL (Verilog, VHDL, SystemVerilog, Chisel). These other options do not have full support and documentation yet. It is possible to use them, but they will require a bigger integration effort.



CSEE-4868: System-on-Chip Platforms

- *Foundation course on the programming, design, and validation of SoCs with emphasis on high-performance embedded applications*
- Offered at Columbia since 2011, moved to upper-level curriculum in Fall 2016
 - required course for CE BS program, elective for MS programs in CS and EE
- Course Goals
 - mastering the HW and SW aspects of integrating heterogeneous components into a complete system
 - designing new components that are reusable across different systems, product generations, and implementation platforms
 - evaluating designs in a multi-objective optimization space

[L. P. Carloni et al. Teaching Heterogeneous Computing with System-Level Design Methods, WCAE 2019]



Upcoming Events: The First OSCAR Workshop!

Open-Source Computer Architecture Research (OSCAR)

Saturday, June 18, 2022 – New York City (co-located with ISCA 2022)



Call for abstracts



Program



Invited Speakers



Venue &
COVID19 info

<https://oscar-workshop.github.io/>

Welcome to OSCAR 2022!

OSCAR 2022 is the first edition of a new workshop on open-source hardware which addresses the wide variety of challenges encountered by both hardware and software engineers in dealing with the increasing heterogeneity of next-generation computer architectures. By providing a venue which brings together researchers from academia, industry and government labs, OSCAR promotes a collaborative approach to foster the efforts of the open-source hardware community in this direction.

In Summary: ESP for Open-Source Hardware

- We contribute **ESP** to the OSH community in order to support the realization of
 - more scalable architectures for SoCs that integrate
 - more heterogeneous components, thanks to a
 - more flexible design methodology, which accommodates different specification languages and design flows
- **ESP** was conceived as a heterogeneous integration platform from the start and tested through years of teaching at Columbia University
- We invite you to **use ESP** for your projects and to **contribute to ESP!**



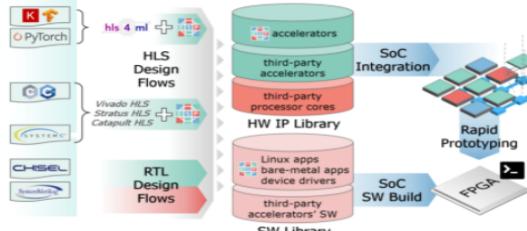
Home Release Resources News Press Team Contact

ESP <https://www.esp.cs.columbia.edu>
the open-source SoC platform



The ESP Vision

ESP is an open-source research platform for heterogeneous system-on-chip design that combines a scalable tile-based architecture and a flexible system-level design methodology.



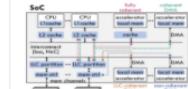
ESP provides three accelerator flows: RTL, high-level synthesis (HLS), machine learning frameworks. All three design flows converge to the ESP automated SoC Integration flow that generates the necessary hardware and software interfaces to rapidly enable full-system prototyping on FPGA.

Overview



- Architecture
 - Tile-based: processor, accelerator, memory, scratchpad and I/O tiles
 - NoC-based
 - Available processors
 - 64-bit Ariane (RISC-V)
 - 32-bit Leon3 (Spac)
 - 32-bit Ibex (RISC-V)
- Accelerator design and integration flows
 - ESP accelerator design flows
 - SystemC with Cadence Stratus HLS
 - C/C++ with Mentor Catapult HLS
 - C/C++ with Xilinx Vivado HLS

Latest Posts



Paper accepted at MICRO 2021

Our paper "Cohesion: Learning-Based Orchestration of Accelerator Coherence in Heterogeneous SoCs" will be presented in October at MICRO 2021.

[Read more](#)

Published: Sep 14, 2021



Paper published in the IEEE Micro special issue on FPGA Computing

Our paper "Accelerator Integration for Open-Source SoC Design" has been published in the IEEE Micro special issue on FPGA Computing (Jul-Aug 2021).

[Read more](#)



Thank you from the **ESP** team!

<https://esp.cs.columbia.edu>

<https://github.com/sld-columbia/esp>



System Level Design Group



COMPUTER SCIENCE

